



Virtual CSSERG Version 1.0

Bee Boring Vanilla · Virtual CSSERG · September 14, 2026

[Full report](#) · [Executive summary](#)

Finding

Virtual CSSERG has a working static publication layer, complete public Project and Scholar indexes, current project-memory files, a guarded deployment component, and a Scholar runner whose documented wiring passes inspection. All seven automated promise groups pass. This is not a completion declaration: several essential manual gates remain.

All three Projects now supply linked one-figure Executive Summaries, Quarto Full Reports, and two-column short PDFs. Predict the Self closed the final automated gap with a two-chapter book, guarded publisher, reciprocal links, and project-specific checks.

Question and method

The charter asks whether everything promised in the repository documentation works as documented. The audit translates those promises into observable evidence, automates only the deterministic and safe subset, and records browser-, credential-, workflow-, and judgment-dependent claims as manual gates.

The standard-library verifier parses public HTML, resolves local resources and fragments, inspects project memory and report artifacts, checks runner syntax and required commands, and exercises deployment with mocked process calls and non-secret placeholders. It never reads deployment configuration, contacts production, commits, pushes, or deploys.

Automated evidence

Repository guidance passes: the governing files, growth procedure, and scaffold command exist. Project memory passes: every Project and the template have PROJECT.md, STATE.md, and DIALOG.md plus valid title, lifecycle status, and substantive-update metadata. Static HTML/CSS passes structural, branding, local-link, Bootstrap, and footer checks. Public catalogs pass complete Project and initial-Scholar coverage, charter-biography fidelity, and metadata-derived update order. Runner wiring and guarded exact-mirror deployment, including its post-transfer checksum/inventory dry run, pass static and mocked checks.

The three-report group now passes. The overall result remains a regression report, not a research-quality score or Version 1 certification.

Selected public designs

The Project Executive Summary now uses the PI-selected evidence brief: question and status, exactly one dense promise

map, then five linked findings and report choices. The Scholar directory uses the selected portrait roster: equal cards, monograms, short introductions, and direct profile links. Authoritative biographies remain on the profiles.

Primary headers contain internal Home, Projects, and Scholars navigation. Every footer places Dr. Jason Jeffrey Jones and CSSERG under About, and GitHub and CC BY 4.0 under Open work. Quarto HTML books remain the default Full Report because their source/output separation, navigation, citations, and growth path suit a static research publication.

Growth workflow

Projects use one of six lifecycle states: Proposed, Active, Blocked, Paused, Completed, or Archived. Publication readiness is separate. STATE.md metadata provides the public title, lifecycle state, and newest substantive update time; the Projects index must match it.

The new create_project.py command validates a permanent lowercase hyphenated slug, refuses overwrite, stages a complete template copy, personalizes its title and state, and renames it atomically. It never publishes an empty Project or displaces PI authorship of the charter. Unit tests cover the success and failure paths.

Scholar creation stays review-led. The PI supplies the name, slug, biography, and assignment. `scholars.json` is the operational source for names, permanent slugs, monograms, and current Projects; a validator rejects duplicates, malformed records, and unknown Projects. The public profiles are checked against it, while PI-authored text remains authoritative for biography. Monograms remain the default; any generated illustration requires repository provenance, accurate alt text, non-photographic presentation, and PI selection.

Manual gates

After the normal September 13 publication, a September 14 public probe found all 109 expected production files byte-identical. That closes the expected-byte check but cannot find remote-only files. Deployment now refuses success when an authenticated post-transfer checksum/inventory dry run reports a missing, changed, or extra path; this new phase awaits observation in the next normal workflow. End-to-end automation still requires witnessed success and failure behavior. Rendered QA requires representative desktop and phone inspection with keyboard and assistive technology; no local browser executable was available. Finally, structural checks cannot establish research validity or substantive completeness.

The approved dialog redesign also remains coordinated future work: one immutable timestamped file per iteration plus a newest-first index, changed together across the orientation, template, runner prompt, active Projects, and verifier.

References

Quarto. (n.d.). *Creating a book*. Retrieved September 11, 2026, from <https://quarto.org/docs/books/>.

Quarto. (n.d.). *HTML accessibility checks*. Retrieved September 11, 2026, from <https://quarto.org/docs/output-formats/html-accessibility.html>.